

Install Kubernetes on Ubuntu 24.04

1) Set Host Name and Update hosts file

SSH to each Ubuntu 24.04 instance and set their respective hostname using hostnamectl command

```
$ sudo hostnamectl set-hostname "k8s-master-noble" // Master Node
```

```
$ sudo hostnamectl set-hostname "k8s-worker01-noble" // Worker Node 1
```

```
$ sudo hostnamectl set-hostname "k8s-worker02-noble" // Worker Node 2
```

Add the following lines to **/etc/hosts** file on each instance.

```
192.168.1.120 k8s-master-noble
```

```
192.168.1.121 k8s-worker01-noble
```

```
192.168.1.122 k8s-worker02-noble
```

2) Disable Swap and Load Kernel Modules

It is highly recommended to disable swap space on your Ubuntu instances so that Kubernetes cluster works smoothly. Run beneath command on each instance to disable swap space.

```
$ sudo swapoff -a
```

```
$ sudo sed -i 's/^(\s*)$/#\1/g' /etc/fstab
```

Now, load the following kernel modules using modprobe command.

```
$ sudo modprobe overlay
```

```
$ sudo modprobe br_netfilter
```

For the permanent loading of these modules, create the file with following content.

```
$ sudo tee /etc/modules-load.d/k8s.conf <<EOF
```

```
overlay
```

```
br_netfilter
```

```
EOF
```

Next, add the kernel parameters like IP forwarding. Create a file and load the parameters using sysctl command,

```
$ sudo tee /etc/sysctl.d/kubernetes.conf <<EOT
```

```
net.bridge.bridge-nf-call-ip6tables = 1
```

```
net.bridge.bridge-nf-call-iptables = 1
```

```
net.ipv4.ip_forward = 1
```

```
EOT
```

To load the above kernel parameters, run

```
$ sudo sysctl --system
```

3) Install and Configure Containerd

Containerd provides the container run time for Kubernetes. So, Install containerd on all three instances.

First install containerd dependencies,

```
$ sudo apt install -y curl gnupg2 software-properties-common apt-transport-https ca-certificates
```

Next, add containerd repository using following set of commands.

```
$ sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo gpg --dearmor -o /etc/apt/trusted.gpg.d/containerd.gpg
```

```
$ sudo add-apt-repository "deb [arch=amd64] https://download.docker.com/linux/ubuntu $(lsb_release -cs) stable"
```

Now, install containerd using following apt command.

```
$ sudo apt update && sudo apt install containerd.io -y
```

Next, configure containerd so that it starts using **SystemdCgroup**. Run beneath commands.

```
$ containerd config default | sudo tee /etc/containerd/config.toml >/dev/null 2>&1
```

```
$ sudo sed -i 's/SystemdCgroup \= false/SystemdCgroup \= true/g' /etc/containerd/config.toml
```

Restart containerd service so that above changes come into the affect.

```
$ sudo systemctl restart containerd
```

4) Add Kubernetes Package Repository

Kubernetes packages are not available in the default package repositories of Ubuntu 24.04, so for its installation first add it's repository. Run these steps on each instance.

Note: At the time of writing this post, latest version of Kubernetes was **1.30**. So you can this version according your requirement.

Download the public signing key for the Kubernetes package repository using curl command.

```
$ curl -fsSL https://pkgs.k8s.io/core:/stable:/v1.30/deb/Release.key | sudo gpg --dearmor -o /etc/apt/keyrings/k8s.gpg
```

Next, add the Kubernetes repository by running following command.

```
$ echo 'deb [signed-by=/etc/apt/keyrings/k8s.gpg] https://pkgs.k8s.io/core:/stable:/v1.30/deb/ /' | sudo tee /etc/apt/sources.list.d/k8s.list
```

5) Install Kubernetes Components (Kubeadm, kubelet & kubectl)

Install Kubernetes components like **Kubeadm**, **kubelet** and **kubectl**, run following apt commands on all the instances.

```
$ sudo apt update
```

```
$ sudo apt install kubelet kubeadm kubectl -y
```

6) Initialize Kubernetes Cluster

As all the prerequisites are met, now we are good to start the installation of Kubernetes on Ubuntu 24.04.

Run following **Kubeadm** command from the master node only to initialize the Kubernetes cluster.

```
$ sudo kubeadm init --control-plane-endpoint=k8s-master-noble  
kubectl config set-cluster kubernetes --server=https://<CIDR>:6443
```

This command will pull the required images for your Kubernetes cluster. Once this command is executed successfully, we

will get the output something like below:

```

[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

You can now join any number of control-plane nodes by copying certificate authorities
and service account keys on each node and then running the following as root:

kubeadm join k8s-master-noble:6443 --token p3sdpk.zn0s060af0089ioa \
--discovery-token-ca-cert-hash sha256:afa3d90b6cd8c5889fca12ea3e9b50659b933ab6c808e2906fd63bde5e695bfd \
--control-plane

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join k8s-master-noble:6443 --token p3sdpk.zn0s060af0089ioa \
--discovery-token-ca-cert-hash sha256:afa3d90b6cd8c5889fca12ea3e9b50659b933ab6c808e2906fd63bde5e695bfd
linuxtech@k8s-master-noble:~$

```

In the output above, we will get a series of commands like how to start interacting with your Kubernetes cluster and command to join any worker node to join this cluster.

On the master node, run following set of commands.

```
$ mkdir -p $HOME/.kube
```

```
$ sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
$ sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Next copy the command to join any worker node from the above output, run it on both the worker nodes. In my case, command would be:

```
$ sudo kubeadm join k8s-master-noble:6443 --token p3sdpk.zn0s060af0089ioa \
--discovery-token-ca-cert-hash sha256:afa3d90b6cd8c5889fca12ea3e9b50659b933ab6c808e2906fd63bde5e695bfd
```

Output from first worker node

```

linuxtech@k8s-worker01-noble:~$
linuxtech@k8s-worker01-noble:~$ sudo kubeadm join k8s-master-noble:6443 --token p3sdpk.zn0s060af0089ioa \
--discovery-token-ca-cert-hash sha256:afa3d90b6cd8c5889fca12ea3e9b50659b933ab6c808e2906fd63bde5e695bfd
[sudo] password for linuxtech:
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 522.838325ms
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

linuxtech@k8s-worker01-noble:~$

```

Similarly output from the second worker node

```

linuxtech@k8s-worker02-noble:~$ sudo kubeadm join k8s-master-noble:6443 --token p3sdpk.zn0s060af0089ioa \
--discovery-token-ca-cert-hash sha256:afa3d90b6cd8c5889fca12ea3e9b50659b933ab6c808e2906fd63bde5e695bfd
[sudo] password for linuxtech:
[preflight] Running pre-flight checks
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-check] Waiting for a healthy kubelet. This can take up to 4m0s
[kubelet-check] The kubelet is healthy after 1.00456183s
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap

This node has joined the cluster:
* Certificate signing request was sent to apiserer and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

linuxtech@k8s-worker02-noble:~$

```

Now head back to the master node and run **kubectl get nodes** command to verify the status of worker nodes.

\$ kubectl get nodes

```

linuxtech@k8s-master-noble:~$
linuxtech@k8s-master-noble:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
k8s-master-noble    NotReady control-plane   16m   v1.30.2
k8s-worker01-noble  NotReady <none>        5m35s v1.30.2
k8s-worker02-noble  NotReady <none>        5m23s v1.30.2
linuxtech@k8s-master-noble:~$

```

7) Install Calico Network Add-on Plugin

To install calico network plugin, run beneath command from the master node only.

\$ kubectl apply -f <https://raw.githubusercontent.com/projectcalico/calico/v3.29.1/manifests/calico.yaml>

```

linuxtech@k8s-master-noble:~$
linuxtech@k8s-master-noble:~$ kubectl apply -f https://raw.githubusercontent.com/projectcalico/calico/v3.28.0/manifests/calico.yaml
poddisruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
serviceaccount/calico-cni-plugin created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgpfilters.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrole.rbac.authorization.k8s.io/calico-cni-plugin created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-cni-plugin created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
linuxtech@k8s-master-noble:~$

```

After the successful installation of calico, nodes status will change to Ready in a minute or two.

\$ kubectl get pods -n kube-system


```
linuxtech@k8s-master-noble:~$ kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
calico-kube-controllers-564985c589-bvwk2	1/1	Running	0	3m1s
calico-node-7t2l8	1/1	Running	0	3m1s
calico-node-gqzlx	1/1	Running	0	3m1s
calico-node-k5t6z	1/1	Running	0	3m1s
coredns-7db6d8ff4d-6qxmm	1/1	Running	0	32m
coredns-7db6d8ff4d-x6p82	1/1	Running	0	32m
etcd-k8s-master-noble	1/1	Running	0	32m
kube-apiserver-k8s-master-noble	1/1	Running	0	32m
kube-controller-manager-k8s-master-noble	1/1	Running	0	32m
kube-proxy-94ld8	1/1	Running	0	22m
kube-proxy-9j76r	1/1	Running	0	21m
kube-proxy-ttdqt	1/1	Running	0	32m
kube-scheduler-k8s-master-noble	1/1	Running	0	32m

```
linuxtech@k8s-master-noble:~$
```

\$ kubectl get nodes

```
linuxtech@k8s-master-noble:~$ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
k8s-master-noble	Ready	control-plane	34m	v1.30.2
k8s-worker01-noble	Ready	<none>	24m	v1.30.2
k8s-worker02-noble	Ready	<none>	23m	v1.30.2

```
linuxtech@k8s-master-noble:~$
```

Output above confirms that nodes are in Ready state.

8) Test Kubernetes Installation

To test the Kubernetes installation, let's create nginx based deployment with replica count 2. Execute the following **kubectl** command from the master node.

\$ kubectl create ns demo-app

\$ kubectl create deployment nginx-app --image nginx --replicas 2 --namespace demo-app

\$ kubectl get deployment -n demo-app

\$ kubectl get pods -n demo-app

```
linuxtech@k8s-master-noble:~$ kubectl create ns demo-app
namespace/demo-app created
linuxtech@k8s-master-noble:~$ kubectl create deployment nginx-app --image nginx --replicas 2 --namespace demo-app
deployment.apps/nginx-app created
linuxtech@k8s-master-noble:~$ kubectl get deployment -n demo-app
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
nginx-app	2/2	2	2	59s

```
linuxtech@k8s-master-noble:~$ kubectl get pods -n demo-app
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-app-69999bf9b8-dwkm8	1/1	Running	0	71s
nginx-app-69999bf9b8-z2wvb	1/1	Running	0	71s

```
linuxtech@k8s-master-noble:~$
```

Next expose this deployment using NodePort type, run

```
$ kubectl expose deployment nginx-app -n demo-app --type NodePort --port 80
```

```
$ kubectl get svc -n demo-app
```

```
linuxtech@k8s-master-noble:~$  
linuxtech@k8s-master-noble:~$ kubectl expose deployment nginx-app -n demo-app --type NodePort --port 80  
service/nginx-app exposed  
linuxtech@k8s-master-noble:~$ kubectl get svc -n demo-app  
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE  
nginx-app     NodePort    10.107.38.103 <none>         80:30336/TCP    24s  
linuxtech@k8s-master-noble:~$  
linuxtech@k8s-master-noble:~$  
linuxtech@k8s-master-noble:~$
```

Now try to access your application using **nodeport** as shown below

```
$ curl http://<Any-worker-IP>:30336
```

If want to check through browser

<http://<Any-worker-IP>:30336>

```
22/06/2024 14:34.06 /home/mobaxterm curl http://192.168.1.121:30336  
<!DOCTYPE html>  
<html>  
<head>  
<title>Welcome to nginx!</title>  
<style>  
html { color-scheme: light dark; }  
body { width: 35em; margin: 0 auto;  
font-family: Tahoma, Verdana, Arial, sans-serif; }  
</style>  
</head>  
<body>  
<h1>Welcome to nginx!</h1>  
<p>If you see this page, the nginx web server is successfully installed and  
working. Further configuration is required.</p>  
  
<p>For online documentation and support please refer to  
<a href="http://nginx.org/">nginx.org</a>.<br/>  
Commercial support is available at  
<a href="http://nginx.com/">nginx.com</a>.</p>  
  
<p><em>Thank you for using nginx.</em></p>  
</body>  
</html>
```

Great, output above confirms that we can access nginx based application outside of our Kubernetes cluster using the nodeport. This confirms that Kubernetes installation is successful.